

An aerial, high-angle view of a city at night. The buildings are dark, but the windows and streets are filled with a dense pattern of warm, yellow-orange lights, creating a bokeh effect. The perspective is looking down from a high vantage point, showing the geometric shapes of the urban landscape.

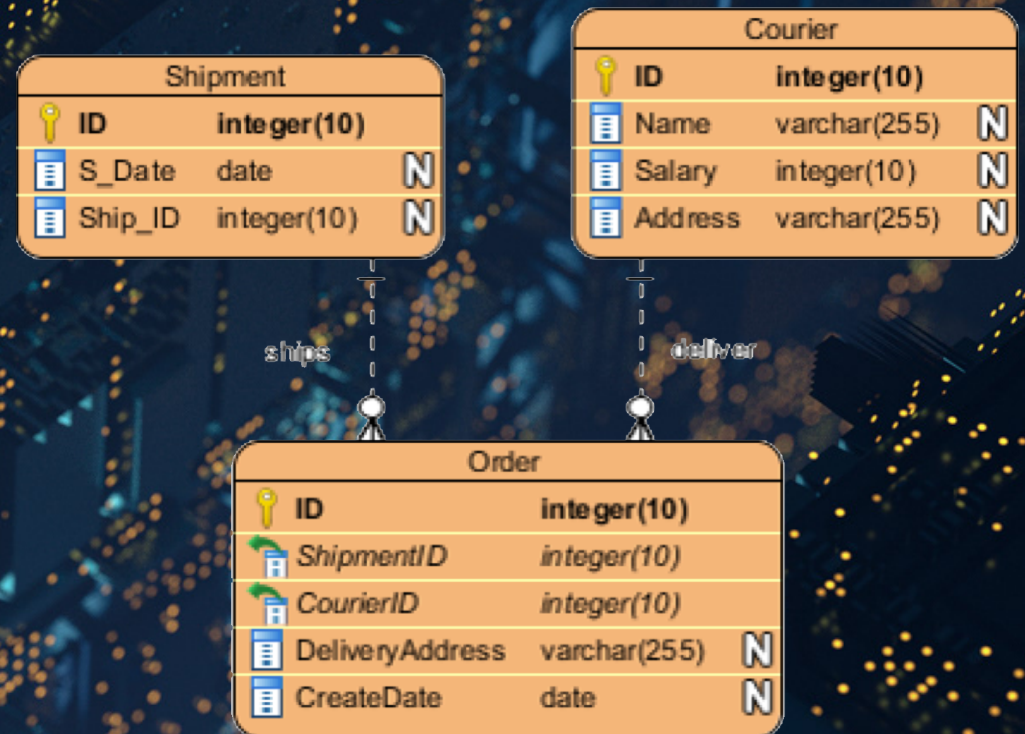
SQL-Joins

Contents

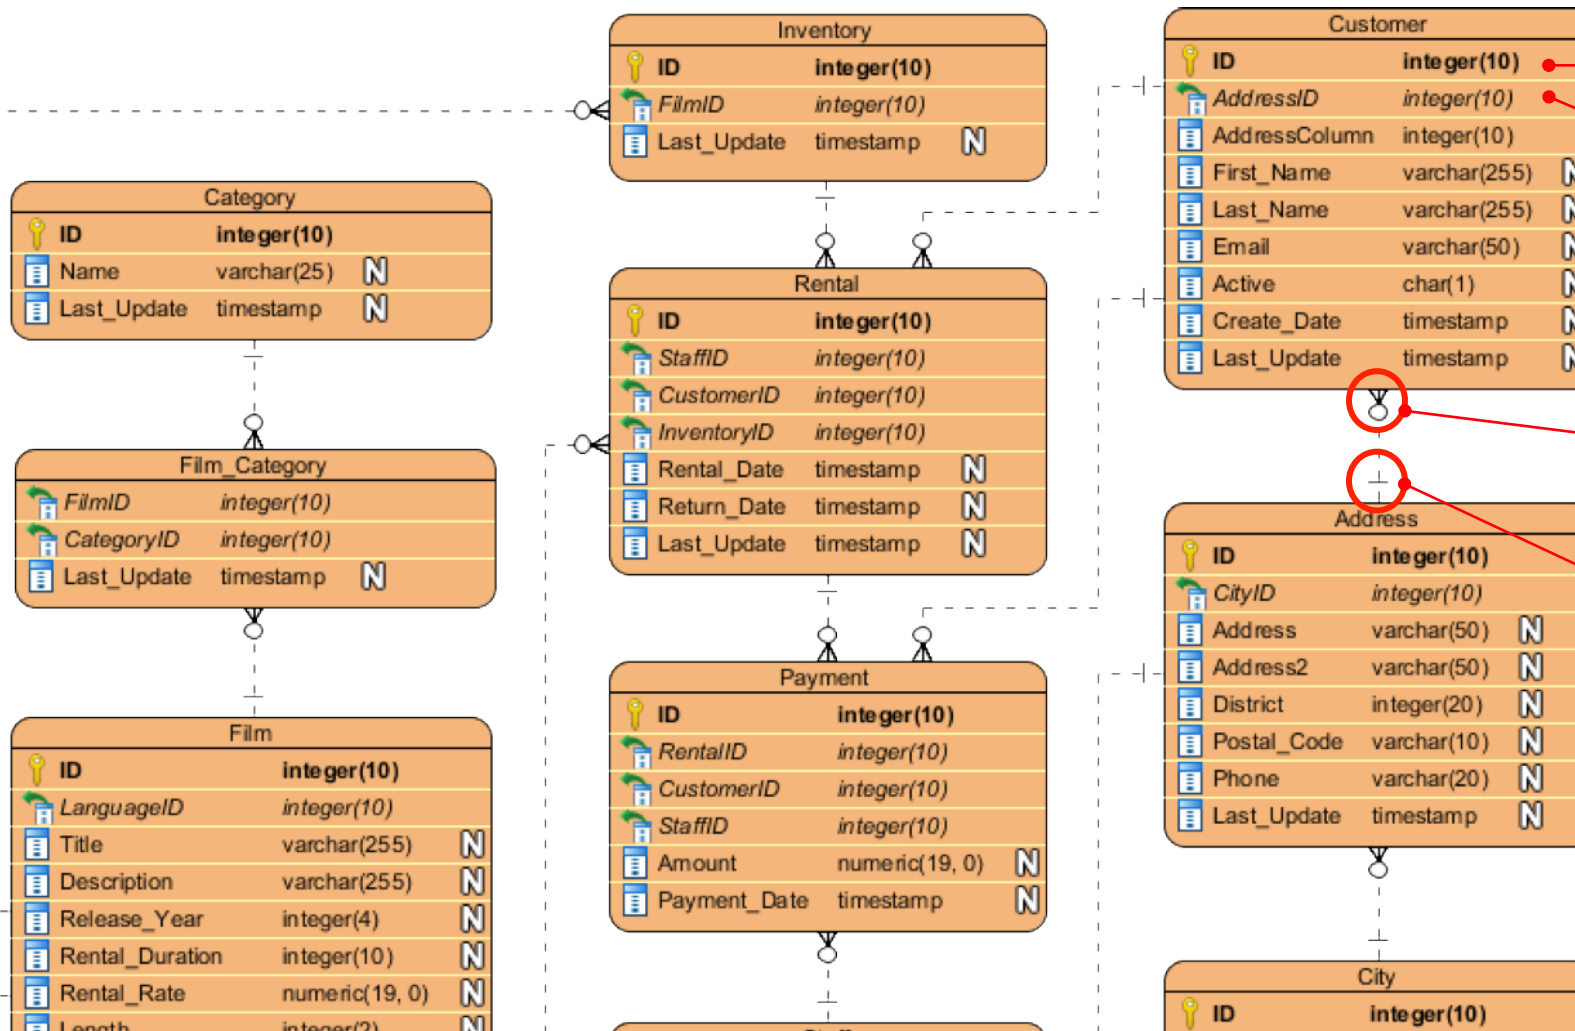
- ERD - Brief Review
- Why using JOIN?
- JOIN Statement
 - INNER JOIN
 - LEFT JOIN
 - RIGHT JOIN
 - FULL OUTER JOIN

Using multiple tables

- Take a look at our friend, the ERD:
- So far, we just used one entity
- Power of SQL comes from the fact we can run queries against multiple tables at once
- JOINS are the statement we use to 'connect' the tables together
- Joins work thanks to **Normalization**, a design technique that reduces data redundancy



EDR Brief Review



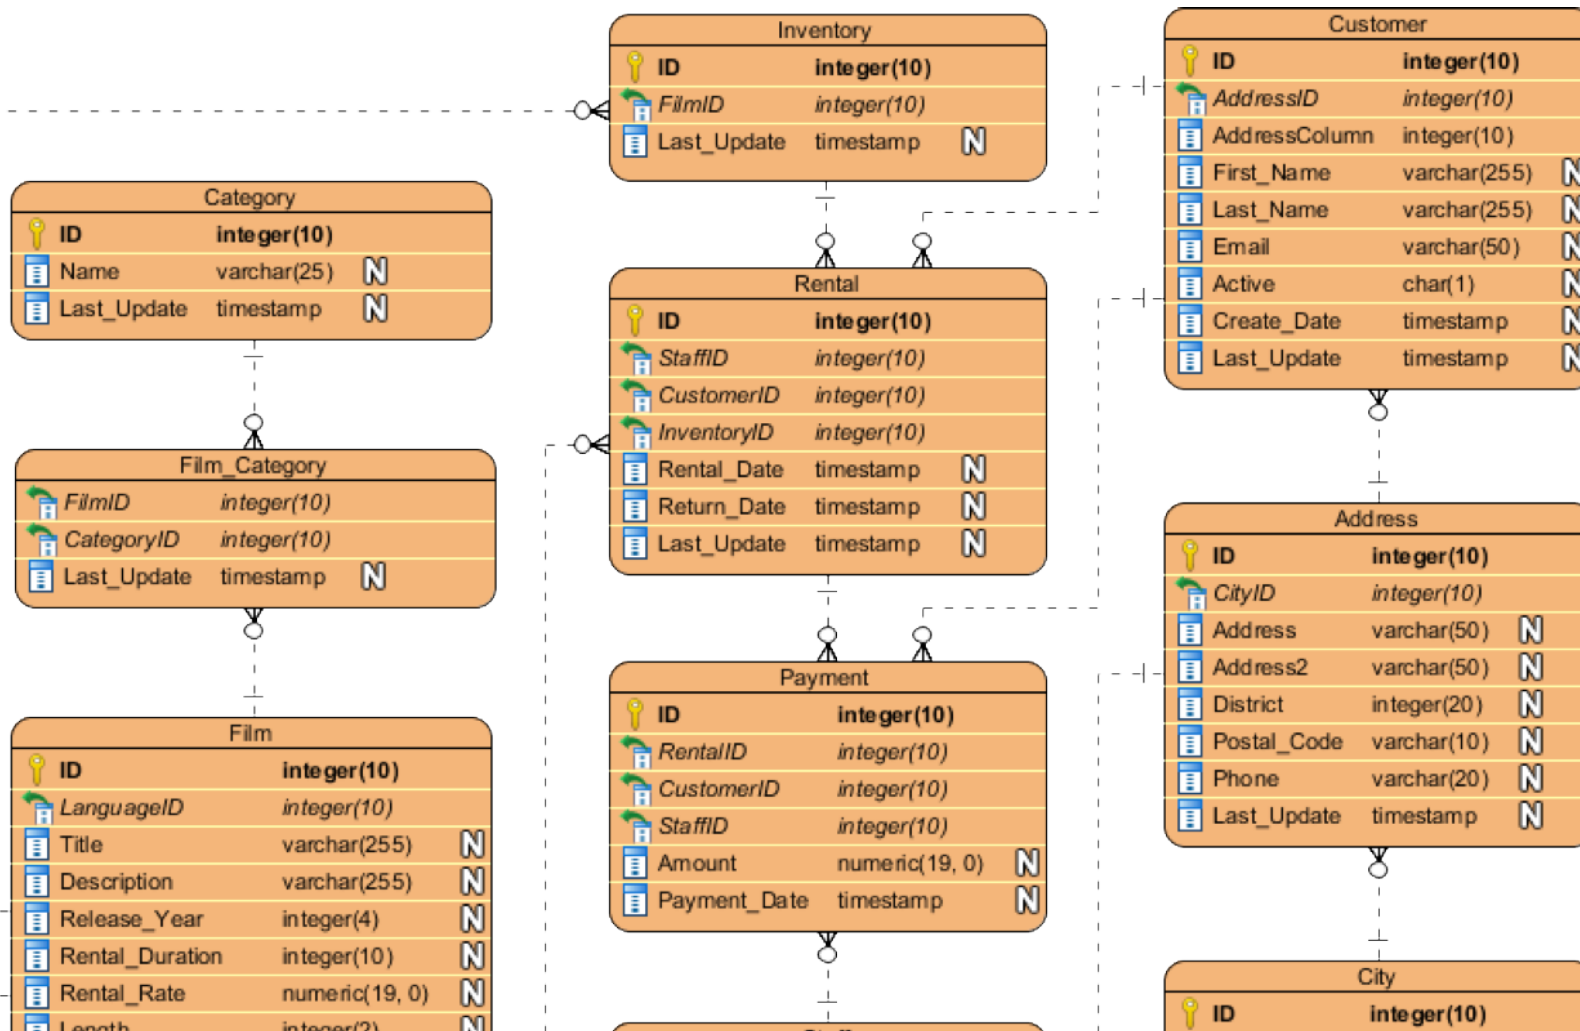
Primary Key (PK):
Unique identifiers

Foreign Key (FK):
Is a column in one table which is a PK

Crow's foot:
Indicates many of the instances of a FK

Dash line:
indicates only one instance of PK will be present in this table

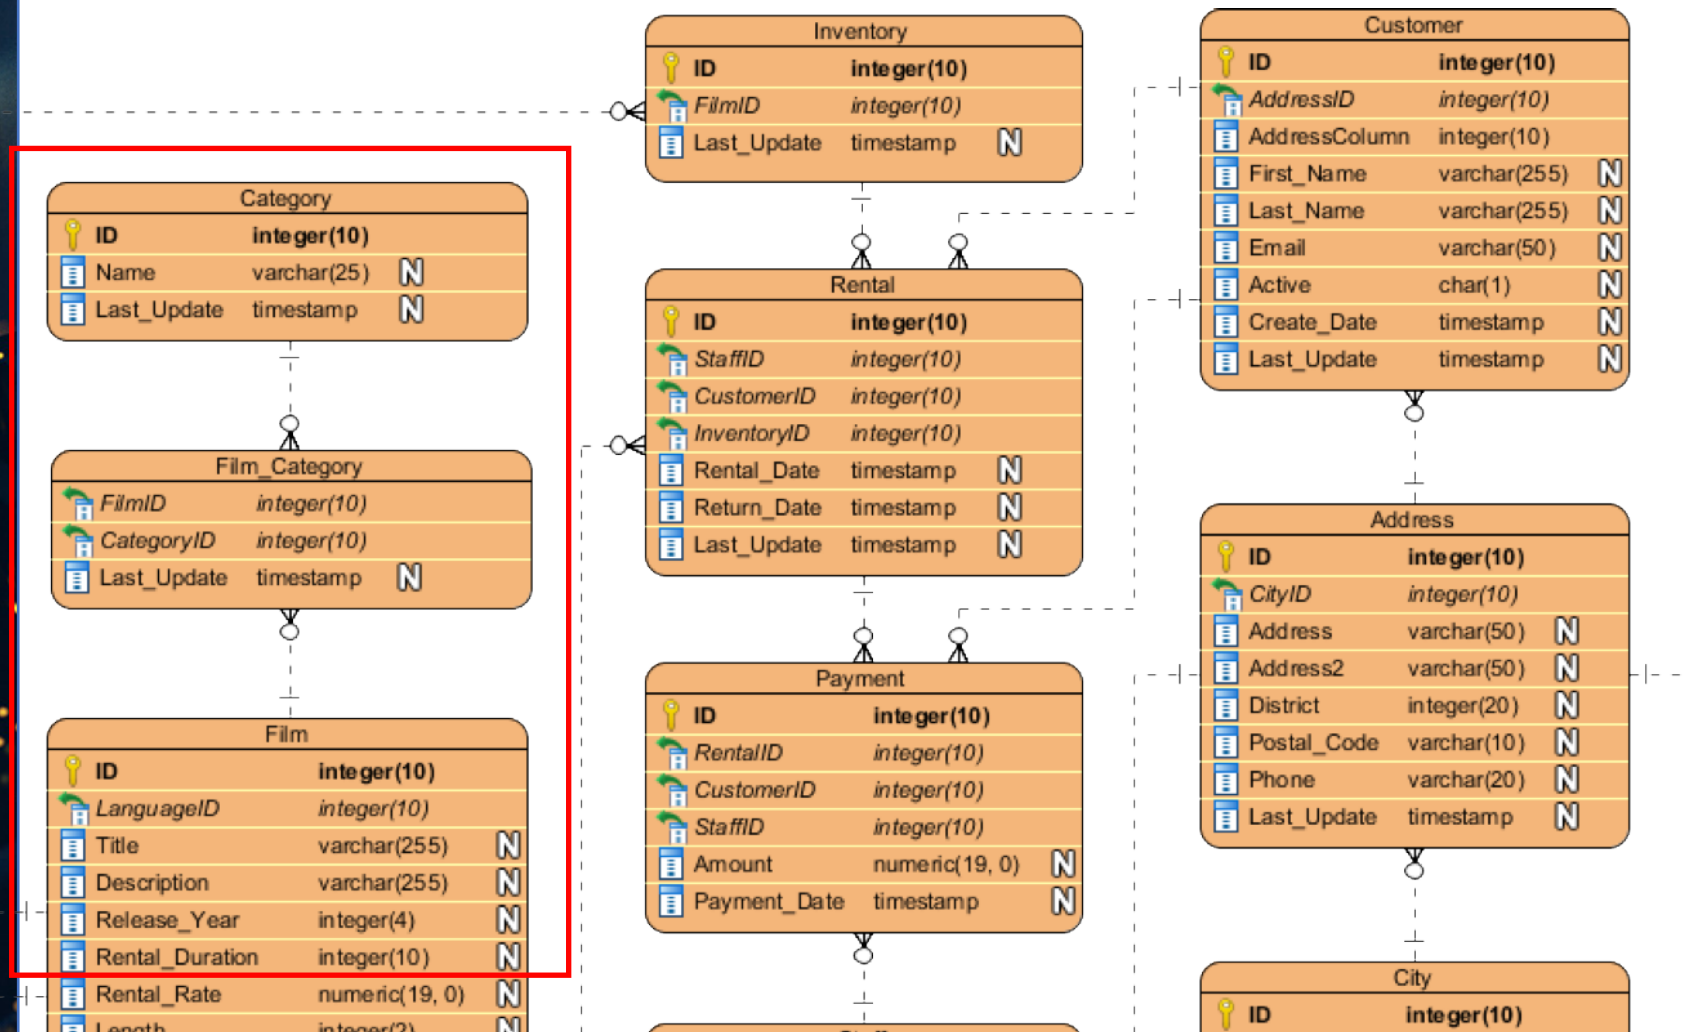
EDR Brief Review



One to Many relationships

- Each Customer has a unique ID.
- ID is the PK of the customer table
- Thus, we will find a single ID for each customer
- This ID is also present in the Rental, Payment
- In those tables, it is present as a FK
- Thus, we might find the same Customer ID for many rental operations.

EDR Brief Review



Many-to-Many relationships

- There are many categories
- There are many films
- One film can have many categories
- One category can be related to many films
- Since PKs have to be unique, many to many relationships have to resort to a intermediate step (*Film Category*)

EDR Brief Review



How can we leverage these relationships?

- As many tables are connected, we can perform interesting analyses
- We can see, for example, the stores in UK

Let's see this without using JOIN first

Why using JOIN

- First, let's find the UK id

```
SELECT * FROM country  
WHERE country = 'United Kingdom';
```

	country_id [PK] integer	country character varying (50)	last_update timestamp without time zone
1	102	United Kingdom	2006-02-15 09:44:00

We can look cities whose country_id = 102

Why using JOIN

```
SELECT * FROM city  
WHERE country_id = 102;
```

	city_id [PK] integer	city character varying (50)	country_id smallint	last_update timestamp without time zone
1	88	Bradford	102	2006-02-15 09:45:25
2	149	Dundee	102	2006-02-15 09:45:25
3	312	London	102	2006-02-15 09:45:25
4	494	Southampton	102	2006-02-15 09:45:25
5	495	Southend-on-Sea	102	2006-02-15 09:45:25
6	496	Southport	102	2006-02-15 09:45:25
7	500	Stockport	102	2006-02-15 09:45:25
8	589	York	102	2006-02-15 09:45:25

Why using JOIN

```
SELECT * FROM address  
WHERE city_id IN (88, 149, 312, 494, 495, 496, 500,  
589);
```

	address_id [PK] integer	address character varying (50)	address2 character varying (50)	district character varying (20)	city_id smallint
1	20	360 Toulouse Parkway		England	495
2	89	1557 Ktahya Boulevard		England	88
3	146	483 Ljubertsy Parkway		Scotland	149
4	256	1497 Yuzhou Drive		England	312
5	482	808 Naala-Porto Parkway		England	500
6	502	1515 Korla Way		England	589
7	517	548 Uruapan Street		Ontario	312
8	562	869 Shikarpur Way		England	496
9	589	1584 Ljubertsy Lane		England	494


Why using JOIN

- The process is tedious
- We have to code the IDs manually
- We can't guarantee that we got all the IDs
- JOIN makes this process simple
- There are four type of JOIN

Why using JOIN


- There are four type of JOIN:

INNER JOIN




Select all records from Table A and Table B, where the join condition is met.

LEFT JOIN




Select all records from Table A, along with records from Table B for which the join condition is met (if at all).

RIGHT JOIN



Select all records from Table B, along with records from Table A for which the join condition is met (if at all).

FULL OUTER JOIN



Select all records from Table A and Table B, regardless of whether the join condition is met or not.

JOIN Statement - INNER

- INNER JOIN: Connects rows based on a condition known as the *join predicate*. It has the same functionality as JOIN (pure syntactic sugar when you have different type of JOINS)

```
SELECT {columns}
FROM {table_1}
INNER JOIN {table_2}
ON {table_1}.{common_key_1} = {table_2}.{common_key_2}
```


JOIN Statement - INNER

```
SELECT *  
FROM country  
JOIN city  
ON country.country_id = city.country_id  
WHERE country = 'United Kingdom';
```


JOIN Statement - INNER

- We can JOIN ON more than one table

```
SELECT *  
FROM address  
JOIN city  
ON address.city_id = city.city_id  
JOIN country  
ON city.country_id = country.country_id  
WHERE country = 'United Kingdom';
```


JOIN Statement - INNER

- We might want to get only the columns that belong to a certain table

```
SELECT address.*  
FROM address  
JOIN city  
ON address.city_id = city.city_id  
JOIN country  
ON city.country_id = country.country_id  
WHERE country = 'United Kingdom';
```


Aliasing

- Aliasing allows us to create temporary variables which we can reference in our query
 - Typically we alias our tables as just the first letter of the table name
 - We can actually omit the AS, but it's in this example for clarity

```
SELECT address, city
FROM address AS ad
JOIN city AS ci
ON ad.city_id = ci.city_id
JOIN country AS co
ON ci.country_id = co.country_id
WHERE country = 'United Kingdom';
```



Practical – Part I

Go to the portal and complete the first practical in this lesson: “Using JOIN statements”

JOIN Statement - LEFT


- There are four type of JOIN:

INNER JOIN




Select all records from Table A and Table B, where the join condition is met.

LEFT JOIN




Select all records from Table A, along with records from Table B for which the join condition is met (if at all).

RIGHT JOIN



Select all records from Table B, along with records from Table A for which the join condition is met (if at all).

FULL OUTER JOIN



Select all records from Table A and Table B, regardless of whether the join condition is met or not.

JOIN Statement - LEFT

LEFT JOIN

1. The left table will have every row returned
2. Matches every row to the row in the right table (based on the ON condition)
 - A. If the ON condition is True, columns from both tables are combined
 - B. If the ON condition is False, a new row is still added but with a NULL value

JOIN Statement - LEFT

There might be some stores with no clients registered:

```
SELECT ad.address_id, cu.*  
FROM address AS ad  
LEFT JOIN customer AS cu  
ON ad.address_id = cu.address_id  
WHERE cu.customer_id IS NULL;
```


JOIN Statement - LEFT


There might be some stores with no clients registered:

595	601	595	1	TERRENCE	GUNDERSON	TERRENCE.GUNDERSON@s...
596	602	596	1	ENRIQUE	FORSYTHE	ENRIQUE.FORSYTHE@sakil...
597	603	597	1	FREDDIE	DUGGAN	FREDDIE.DUGGAN@sakilac...
598	604	598	1	WADE	DELVALLE	WADE.DELVALLE@sakilacus...
599	605	599	2	AUSTIN	CINTRON	AUSTIN.CINTRON@sakilacu...
600	2	[null]	[null]	[null]	[null]	[null]
601	4	[null]	[null]	[null]	[null]	[null]
602	1	[null]	[null]	[null]	[null]	[null]
603	3	[null]	[null]	[null]	[null]	[null]

JOIN Statement - RIGHT


- There are four type of JOIN:

INNER JOIN




Select all records from Table A and Table B, where the join condition is met.

LEFT JOIN




Select all records from Table A, along with records from Table B for which the join condition is met (if at all).

RIGHT JOIN



Select all records from Table B, along with records from Table A for which the join condition is met (if at all).

FULL OUTER JOIN



Select all records from Table A and Table B, regardless of whether the join condition is met or not.

JOIN Statement - RIGHT

RIGHT JOIN


- Very similar to LEFT JOIN
- The output is similar, but just switched positions

```
SELECT as.address_id, cu.*  
FROM customer AS cu  
RIGHT JOIN address AS ad  
ON cu.address_id = ad.address_id;
```


JOIN Statement - FULL OUTER


- There are four type of JOIN:

INNER JOIN




Select all records from Table A and Table B, where the join condition is met.

LEFT JOIN




Select all records from Table A, along with records from Table B for which the join condition is met (if at all).

RIGHT JOIN



Select all records from Table B, along with records from Table A for which the join condition is met (if at all).

FULL OUTER JOIN



Select all records from Table A and Table B, regardless of whether the join condition is met or not.

JOIN Statement - FULL OUTER

FULL [OUTER] JOIN

- Take all data from both tables, regardless of the matches.
- Let's see the example from the previous slides (The employees exercise)

```
SELECT *  
FROM employee_details AS det  
FULL OUTER JOIN employee_salary AS sal  
ON det.employee_id = sal.employee_id;
```


JOIN Statement - FULL OUTER

	employee_id integer	employee_name character varying (20)	employee_id integer	employee_name character varying (20)	salary integer
1	1	Mr. Pink	1	Mr. Pink	50000
2	2	Mr. Blonde	2	Mr. Blonde	48000
3	3	Mr. Orange	3	Mr. Orange	65000
4	4	Mr. White	[null]	[null]	[null]
5	5	Mr. Brown	[null]	[null]	[null]
6	6	Eddie	6	Eddie	90000
7	7	Joe	7	Joe	120000
8	[null]	[null]	8	Mr Blue	30000

JOIN Statement - FULL OUTER

FULL [OUTER] JOIN

- Rare use case
- Discussion about its uses: <https://stackoverflow.com/questions/2094793/when-is-a-good-situation-to-use-a-full-outer-join>


JOIN Statement

- There are four type of JOIN:



JOIN Statement

`SELECT * FROM a
INNER JOIN b ON a.key = b.key`



`SELECT * FROM a
LEFT JOIN b ON a.key = b.key`



`SELECT * FROM a
RIGHT JOIN b ON a.key = b.key`



`SELECT * FROM a
LEFT JOIN b ON a.key = b.key
WHERE b.key IS NULL`



`SELECT * FROM a
RIGHT JOIN b ON a.key = b.key
WHERE a.key IS NULL`



`SELECT * FROM a
FULL JOIN b ON a.key = b.key`



`SELECT * FROM a
FULL JOIN b ON a.key = b.key
WHERE a.key IS NULL OR b.key IS NULL`



Practical - Part II

Use the following tables to complete the Practical

customerID	customerName
1	Homer
2	Marge
3	Bart
4	Lisa
5	Maggie
6	Moe

orderID	customerID	item
1	1	Beer
2	2	Hair product
3	2	Dress
4	3	Juice
5	3	Magazine
6	6	Peanuts

Practical - Part II

1. Given the following SQL statement:

```
SELECT c.customerID, c.customerName, o.item  
FROM customer AS c  
INNER JOIN order AS o  
ON c.customerID = o.customerID
```

What is the:

- A. Number of columns in the returned table?
- B. Number of rows in the returned table?
- C. Number of times customerID "2" would show up?
- D. Number of times customerID "5" would show up?

Practical - Part II

1. Given the following SQL statement:

```
SELECT c.customerID, c.customerName, o.item  
FROM customer AS c  
LEFT JOIN order AS o  
ON c.customerID = o.customerID
```

What is the:

- A. Number of columns in the returned table?
- B. Number of rows in the returned table?
- C. Number of times customerID "2" would show up?
- D. Number of times customerID "5" would show up?